

# Parquet-MR acceleration using VectorAPI

Jatin Bhateja (Intel)  
Fang Xie (Intel)

Bitpack algorithm is very popular in database and big data storage system like Parquet, ORC. Vector API optimization provide better performance for en/decoding bitpack.


Below are performance improvemement reports

Bit Width	1	2	3	4	5	6	7	8	9	10
Scalar time (ms):	3193	163	145	81	99	99	100	38	91	96
Vector time (ms):	873	28	41	15	25	19	17	8	12	12
Performance gain	3.66x	5.82x	3.53x	5.4x	3.96x	5.21x	5.88x	4.75x	7.58x	8.0x

## Vectorized Execution: bit-packing decode

### Test2: Reader test

- ❖ Location: Spark  
VectorizedParquetRecordReader
- ❖ Read mode: Batch column reader
- ❖ Bit width: 7
- ❖ Compression: false
- ❖ Batch Size: 4096
- ❖ Column: 1~4 columns
- ❖ Rows: 10k~3000k




# Vectorized Execution: **bit-packing decode**

## Test2: Reader test

- ❖ Location: Spark  
VectorizedParquetRecordReader
- ❖ Read mode: Batch column reader
- ❖ Bit width: 7
- ❖ Compression: false
- ❖ Batch Size: 4096
- ❖ Column: 1~4 columns
- ❖ Rows: 3Mills~30Mills


Scalar vs Vector time spend (3~30mill rows)  
lower is better



## Vectorized Execution: bit-packing decode

Test2: Reader test

- ❖ Location: Spark  
VectorizedParquetRecordReader
- ❖ Read mode: Batch column reader
- ❖ Bit width: 7
- ❖ Compression: false
- ❖ Batch Size: 4096
- ❖ Column: 1~4 columns
- ❖ Rows: 30Mills~100Mills



Test sql query in Spark3.3.0 with OpenJDK17, the bitpack optimization improved sql query by 2x~3x in some special case (deep scan file cases)

# Spark SQL query with the optimization of parquet

## Hardware configuration

Architecture: x86\_64  
CPU op-mode(s): 32-bit, 64-bit  
Byte Order: Little Endian  
CPU(s): 112  
On-line CPU(s) list: 0-111  
Thread(s) per core: 2  
Core(s) per socket: 28  
Socket(s): 2  
NUMA node(s): 2  
Vendor ID: GenuineIntel  
BIOS Vendor ID: Intel(R) Corporation  
CPU family: 6  
Model: 106  
Model name: Intel(R) Xeon(R) Platinum 8372C CPU @ 3.20GHz  
BIOS Model name: Intel(R) Xeon(R) Platinum 8372C CPU @ 3.20GHz  
Stepping: 6  
CPU MHz: 3500.000  
CPU max MHz: 3500.0000  
CPU min MHz: 800.0000  
BogoMIPS: 6400.00  
Virtualization: VT-x  
L1d cache: 48K  
L1i cache: 32K  
L2 cache: 1280K  
L3 cache: 43008K  
NUMA node0 CPU(s): 0-27,56-83  
NUMA node1 CPU(s): 28-55,84-111

## Spark configuration

Standalone cluster 1 host  
1 master  
1 worker  
Master memory: 2g  
Executor memory: 32g  
Driver memory: 5g


# Spark SQL query with the optimization of parquet

select count(\*) from ParquetTable where int32\_field>2000 and int32\_field1<2000


## Test1 : SQL test

- ❖ Location: Spark SQL
- ❖ Data type: INT32
- ❖ Compression: Snappy
- ❖ Column: 4 columns
- ❖ Rows: 1000000000
- ❖ File size: 652M
- ❖ Perf gain: 15%

Scalar API v1 snappy (1~6400) for bitpack: "select count(\*) from ParquetTable where int32\_field>2000 and int32\_field1<2000 " everage time:554.0ms



Vector API optimization for v1 snappy (0~6400)bitpack: "select count(\*) from ParquetTable where int32\_field>2000 and int32\_field1<2000 " everage time:469.52ms




# Spark SQL query with the optimization of parquet

select sum(int32\_field2) from ParquetTable where int32\_field1 > 2000 group by int32\_field2

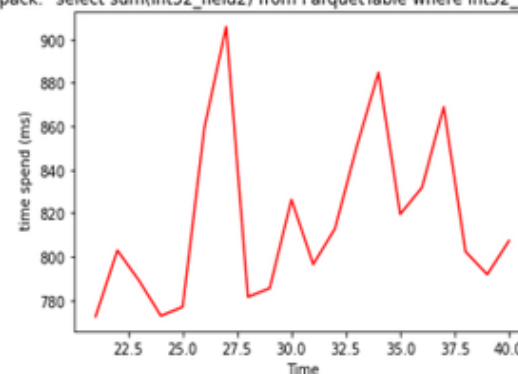
## Test1 : SQL test

- ❖ Location: Spark SQL
- ❖ Data type: INT32
- ❖ Compression: Snappy
- ❖ Column: 4 columns
- ❖ Rows: 1000000000
- ❖ File size: 652M
- ❖ Perf gain: 19%

Scalar API v1 snappy (1~6400) for bitpack: "select sum(int32\_field2) from ParquetTable where int32\_field1 > 2000 group by int32\_field2" average time:1011.31ms



Vector API optimization for v1 snappy (0~6400)bitpack: "select sum(int32\_field2) from ParquetTable where int32\_field1 > 2000 group by int32\_field2" average time:816.92ms




# Spark SQL query with the optimization of parquet

select \* from ParquetTable where int32\_field1 =6399 order by int32\_field2

## Test1 : SQL test

- ❖ Location: Spark SQL
- ❖ Data type: INT32
- ❖ Compression: Snappy
- ❖ Column: 4 columns
- ❖ Rows: 1000000000
- ❖ File size: 652M
- ❖ Perf gain: 20%

Scalar API v1 snappy (1~6400) for bitpack: "select \* from ParquetTable where int32\_field1 =6399 order by int32\_field2" average time:866.38ms



Vector API optimization for v1 snappy (0~6400)bitpack: "select \* from ParquetTable where int32\_field1=6399 order by int32\_field2" average time:675.79ms

